

# Algorithms for data streams: Streaming models, graph streams

Data stream  
models

Graph  
streams

Connectedness

Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window



# The data

## Data stream models

## Graph streams

Connectedness

Maximum matching

## Sampling

Reservoir sampling

Sampling in a sliding window

# The data

- Data arrives as sequence of items.
- Sometimes continuously and at high speed.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# The data

- Data arrives as sequence of items.
- Sometimes continuously and at high speed.
- Can't store them all in main memory.
- Can't read again; or reading again has a cost.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# The data

- Data arrives as sequence of items.
- Sometimes continuously and at high speed.
- Can't store them all in main memory.
- Can't read again; or reading again has a cost.
- We abstract the data to a particular feature, the data field of interest **the label**.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# The data

## Data stream models

## Graph streams

Connectedness

Maximum matching

## Sampling

Reservoir sampling

Sampling in a sliding window

# The data

- We have a set of  $n$  labels  $\Sigma$  and our input is a stream  $s = x_1, x_2, x_3, \dots x_m$ , where each  $x_i \in \Sigma$ .

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# The data

- We have a set of  $n$  labels  $\Sigma$  and our input is a stream  $s = x_1, x_2, x_3, \dots, x_m$ , where each  $x_i \in \Sigma$ .
- Take into account that some times we do not know in advance the length of the stream.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window



# The data

- We have a set of  $n$  labels  $\Sigma$  and our input is a stream  $s = x_1, x_2, x_3, \dots, x_m$ , where each  $x_i \in \Sigma$ .
- Take into account that some times we do not know in advance the length of the stream.
- **Goal** Compute a function of stream, e.g., median, number of distinct elements, longest increasing sequence.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Why has it become popular?

## ■ Practical appeal:

### Data stream models

### Graph streams

Connectedness

Maximum matching

### Sampling

Reservoir sampling

Sampling in a sliding window

# Why has it become popular?

- Practical appeal:

- Faster networks, cheaper data storage, ubiquitous data-logging results in massive amount of data to be processed.
- Applications to network monitoring, query planning, I/O efficiency for massive data, sensor networks aggregation, ...

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Why has it become popular?

## ■ Practical appeal:

- Faster networks, cheaper data storage, ubiquitous data-logging results in massive amount of data to be processed.
- Applications to network monitoring, query planning, I/O efficiency for massive data, sensor networks aggregation, ...

## ■ Theoretical Appeal:

- Easy to state problems but hard to solve.
- Links to communication complexity, compressed sensing, embeddings, pseudo-random generators, approximation, parallel computation, ...

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Why has it become popular?

## ■ Practical appeal:

- Faster networks, cheaper data storage, ubiquitous data-logging results in massive amount of data to be processed.
- Applications to network monitoring, query planning, I/O efficiency for massive data, sensor networks aggregation, ...

## ■ Theoretical Appeal:

- Easy to state problems but hard to solve.
- Links to communication complexity, compressed sensing, embeddings, pseudo-random generators, approximation, parallel computation, ...
- Origins in 70's but has become popular in this century because of growing theory and very applicable.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# The computational data stream models

## ■ Classical streaming model

### Data stream models

#### Graph streams

Connectedness

Maximum matching

#### Sampling

Reservoir sampling

Sampling in a sliding window

# The computational data stream models

## ■ Classical streaming model

- The data stream is accessed sequentially.
- The processing is done sequentially using a small working memory  $O(\text{polylog } n)$ .

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# The computational data stream models

## ■ Classical streaming model

- The data stream is accessed sequentially.
- The processing is done sequentially using a small working memory  $O(\text{polylog } n)$ .
- Measures of complexity: number of **passes** over the data, the size of the working **memory**, the **per-item** processing time.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window



# The computational data stream models

## ■ Classical streaming model

- The data stream is accessed sequentially.
- The processing is done sequentially using a small working memory  $O(\text{polylog } n)$ .
- Measures of complexity: number of **passes** over the data, the size of the working **memory**, the **per-item** processing time.

## ■ Semi-streaming model

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# The computational data stream models

## Data stream models

### Graph streams

Connectedness  
Maximum matching

### Sampling

Reservoir sampling  
Sampling in a sliding window

## ■ Classical streaming model

- The data stream is accessed sequentially.
- The processing is done sequentially using a small working memory  $O(\text{polylog } n)$ .
- Measures of complexity: number of **passes** over the data, the size of the working **memory**, the **per-item** processing time.

## ■ Semi-streaming model

- Usual for graph problems.
- Working memory is  $O(n \text{ polylog } n)$ , for a graph with  $n$  vertices.
- Enough space to store vertices but not for storing all the edges.

# Algorithmic goals

- Data streams are potentially of unbounded size.
- As the amount of computation and memory is limited it might be impossible to provide exact answers.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Algorithmic goals

- Data streams are potentially of unbounded size.
- As the amount of computation and memory is limited it might be impossible to provide exact answers.
- Algorithms use randomization and seek for an approximate answer.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Algorithmic goals

## Data stream models

### Graph streams

Connectedness  
Maximum matching

### Sampling

Reservoir sampling  
Sampling in a sliding window

- Data streams are potentially of unbounded size.
- As the amount of computation and memory is limited it might be impossible to provide exact answers.
- Algorithms use randomization and seek for an approximate answer.
- Typical approach:
  - Build up a **synopsis data structure**
  - It should be enough to compute answers with a high confidence level.

# Streams that describe graphs

- $G$  undirected
- on  $[n]$  vertices
- the stream describes the edges of  $G$
- we assume that an edge appears only once in the stream

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Streams that describe graphs

- $G$  undirected
- on  $[n]$  vertices
- the stream describes the edges of  $G$
- we assume that an edge appears only once in the stream
- We want to keep a DS that allows to answer queries about a graph property
- $O(n \log n)$  memory is reasonable as we are working on the semi-streaming model.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Connectedness

- **Problem:** Decide whether or not the input graph  $G$ , given by a stream, is connected.

Data stream  
models

Graph  
streams

Connectedness

Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window



# Connectedness

- **Problem:** Decide whether or not the input graph  $G$ , given by a stream, is connected.
- **Algorithm:**

Data stream  
models

Graph  
streams

Connectedness

Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Connectedness

- **Problem:** Decide whether or not the input graph  $G$ , given by a stream, is connected.
- **Algorithm:**
  - Maintain a spanning forest  $H$  of the seen graph
  - On query answer according to  $H$

Data stream  
models

Graph  
streams

Connectedness

Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Connectedness

- **Problem:** Decide whether or not the input graph  $G$ , given by a stream, is connected.
- **Algorithm:**
  - Maintain a spanning forest  $H$  of the seen graph
  - On query answer according to  $H$
- $G$  connected iff admits a spanning tree,

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Connectedness

- **Problem:** Decide whether or not the input graph  $G$ , given by a stream, is connected.
- **Algorithm:**
  - Maintain a spanning forest  $H$  of the seen graph
  - On query answer according to  $H$
- $G$  connected iff admits a spanning tree, the algorithm is correct

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Connectedness

- **Problem:** Decide whether or not the input graph  $G$ , given by a stream, is connected.
- **Algorithm:**
  - Maintain a spanning forest  $H$  of the seen graph
  - On query answer according to  $H$
- $G$  connected iff admits a spanning tree, the algorithm is correct
- 1 pass,  $O(n \log n)$  memory,

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Connectedness

- **Problem:** Decide whether or not the input graph  $G$ , given by a stream, is connected.
- **Algorithm:**
  - Maintain a spanning forest  $H$  of the seen graph
  - On query answer according to  $H$
- $G$  connected iff admits a spanning tree, the algorithm is correct
- 1 pass,  $O(n \log n)$  memory, using a union find DS amortized  $O(\alpha(n))$  per item

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Maximum matching

- **Problem:** Find a maximum matching in  $G$ , given by a stream.

Data stream  
models

Graph  
streams

Connectedness

Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Maximum matching

- **Problem:** Find a maximum matching in  $G$ , given by a stream.
- The algorithm maintains a matching  $M$ .

Data stream  
models

Graph  
streams

Connectedness

Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window



# Maximum matching

- **Problem:** Find a maximum matching in  $G$ , given by a stream.
- The algorithm maintains a matching  $M$ .
- **Algorithm:**

Data stream  
models

Graph  
streams

Connectedness

Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Maximum matching

- **Problem:** Find a maximum matching in  $G$ , given by a stream.
- The algorithm maintains a matching  $M$ .
- **Algorithm:**
  - 1: **procedure** MMATCHING(int  $n$ , stream  $s$ , double  $t$ )
  - 2:      $M = \emptyset$
  - 3:     **while** not  $s.end()$  **do**
  - 4:          $(u, v) = s.read()$
  - 5:         **if**  $M \cup (u, v)$  is a matching **then**
  - 6:              $M = M \cup \{(u, v)\}$
  - 7:     On query, report  $M$

Data stream  
models

Graph  
streams

Connectedness

Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Maximum matching: Analysis

Data stream  
models

Graph  
streams

Connectedness

**Maximum matching**

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Maximum matching: Analysis

- 1 pass,  $O(n \log n)$  space,  $O(1)$  ops. per item
- $M$  is a maximal matching that provides an estimation  $\hat{f}$  of the size  $f$  of a maximum matching.
- $\hat{f}$  is a 2 approximation to  $f$ .  
Because, at least one vertex of each edge of  $M$  must be matched by an edge in a maximum matching.

Data stream  
models

Graph  
streams

Connectedness

Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Sampling

Data stream  
models

Graph  
streams

Connectedness

Maximum matching

Sampling

**Reservoir sampling**

Sampling in a sliding  
window

# Sampling

- Sampling is a general technique for tackling massive amounts of data.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

**Reservoir sampling**

Sampling in a sliding  
window

# Sampling

- Sampling is a general technique for tackling massive amounts of data.
- **Example:** To compute the median packet size of some IP packets, we could just sample some and use the median of the sample as an estimate for the true median. Statistical arguments relate the size of the sample to the accuracy of the estimate.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Sampling

- Sampling is a general technique for tackling massive amounts of data.
- **Example:** To compute the median packet size of some IP packets, we could just sample some and use the median of the sample as an estimate for the true median. Statistical arguments relate the size of the sample to the accuracy of the estimate.
- **Challenge:** But how do you take a sample from a stream of unknown length or from a **sliding window**?

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window



# Reservoir Sampling (Vitter 1985)

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

**Reservoir sampling**

Sampling in a sliding  
window

# Reservoir Sampling (Vitter 1985)

- **Problem:** Maintain a uniform sample  $x$  from a stream  $s$  of unknown length.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Reservoir Sampling (Vitter 1985)

- **Problem:** Maintain a uniform sample  $x$  from a stream  $s$  of unknown length.  
The selected item should be any of the seen ones with uniform probability.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Reservoir Sampling (Vitter 1985)

- **Problem:** Maintain a uniform sample  $x$  from a stream  $s$  of unknown length.  
The selected item should be any of the seen ones with uniform probability.
- **Algorithm:**

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Reservoir Sampling (Vitter 1985)

- **Problem:** Maintain a uniform sample  $x$  from a stream  $s$  of unknown length.

The selected item should be any of the seen ones with uniform probability.

- **Algorithm:**
  - Initially  $x = x_1$
  - On seeing the  $t$ -th element,  $x = x_t$  with probability  $1/t$

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Reservoir Sampling (Vitter 1985)

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

- **Problem:** Maintain a uniform sample  $x$  from a stream  $s$  of unknown length.

The selected item should be any of the seen ones with uniform probability.

- **Algorithm:**
  - Initially  $x = x_1$
  - On seeing the  $t$ -th element,  $x = x_t$  with probability  $1/t$
- **Analysis:**

# Reservoir Sampling (Vitter 1985)

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

- **Problem:** Maintain a uniform sample  $x$  from a stream  $s$  of unknown length.  
The selected item should be any of the seen ones with uniform probability.
- **Algorithm:**
  - Initially  $x = x_1$
  - On seeing the  $t$ -th element,  $x = x_t$  with probability  $1/t$
- **Analysis:**
  - 1 pass,  $O(\log n)$  memory (in bits), and  $O(1)$  time (in operations) per item.

# Reservoir Sampling (Vitter 1985)

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

- **Problem:** Maintain a uniform sample  $x$  from a stream  $s$  of unknown length.

The selected item should be any of the seen ones with uniform probability.

- **Algorithm:**
  - Initially  $x = x_1$
  - On seeing the  $t$ -th element,  $x = x_t$  with probability  $1/t$
- **Analysis:**
  - 1 pass,  $O(\log n)$  memory (in bits), and  $O(1)$  time (in operations) per item.
  - **Quality?**  
What is the probability that  $x = x_i$  at some time  $t \geq i$ ?



# Reservoir Sampling: Quality

- At any time step  $t$ , for  $i \leq t$ ,  $Pr[x = x_i] = 1/t$

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

**Reservoir sampling**

Sampling in a sliding  
window

# Reservoir Sampling: Quality

- At any time step  $t$ , for  $i \leq t$ ,  $Pr[x = x_i] = 1/t$
- The proof is by induction on  $t$ .

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Reservoir Sampling: Quality

- At any time step  $t$ , for  $i \leq t$ ,  $Pr[x = x_i] = 1/t$
- The proof is by induction on  $t$ .
  - Base  $t = 1$ :  $Pr[x = x_1] = 1$ .
  - Induction hypothesis: true for time steps up to  $t - 1$ 
    - $Pr[x = x_t] = 1/t$
    - For  $i < t$ ,  $x = x_i$  only when  $x_t$  is not selected and  $x_i$  was the sampled element at step  $t - 1$ . By induction hypothesis we have

$$Pr[x = x_t] = \left(1 - \frac{1}{t}\right) \frac{1}{t-1} = \frac{1}{t}$$

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Reservoir Sampling II

- **Problem:** Maintain a uniform sample  $X$  of size  $k$  from a stream of unknown length.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Reservoir Sampling II

- **Problem:** Maintain a uniform sample  $X$  of size  $k$  from a stream of unknown length.
- **Algorithm:**

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Reservoir Sampling II

- **Problem:** Maintain a uniform sample  $X$  of size  $k$  from a stream of unknown length.
- **Algorithm:**
  - Initially  $X = \{x_1, \dots, x_k\}$ .
  - On seeing the  $t$ -th element,  $t > k$ , select  $x_t$  to be added to  $X$  with probability  $k/t$ .
  - If  $x_t$  is selected to be added, select uniformly at random an element from  $X$ , remove it and add  $x_t$ .

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Reservoir Sampling II

- **Problem:** Maintain a uniform sample  $X$  of size  $k$  from a stream of unknown length.
- **Algorithm:**
  - Initially  $X = \{x_1, \dots, x_k\}$ .
  - On seeing the  $t$ -th element,  $t > k$ , select  $x_t$  to be added to  $X$  with probability  $k/t$ .
  - If  $x_t$  is selected to be added, select uniformly at random an element from  $X$ , remove it and add  $x_t$ .
- **Analysis:**

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Reservoir Sampling II

- **Problem:** Maintain a uniform sample  $X$  of size  $k$  from a stream of unknown length.
- **Algorithm:**
  - Initially  $X = \{x_1, \dots, x_k\}$ .
  - On seeing the  $t$ -th element,  $t > k$ , select  $x_t$  to be added to  $X$  with probability  $k/t$ .
  - If  $x_t$  is selected to be added, select uniformly at random an element from  $X$ , remove it and add  $x_t$ .
- **Analysis:**
  - 1 pass,  $O(k \log n)$  memory, and  $O(1)$  time per item.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window



# Reservoir Sampling II

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

- **Problem:** Maintain a uniform sample  $X$  of size  $k$  from a stream of unknown length.
- **Algorithm:**
  - Initially  $X = \{x_1, \dots, x_k\}$ .
  - On seeing the  $t$ -th element,  $t > k$ , select  $x_t$  to be added to  $X$  with probability  $k/t$ .
  - If  $x_t$  is selected to be added, select uniformly at random an element from  $X$ , remove it and add  $x_t$ .
- **Analysis:**
  - 1 pass,  $O(k \log n)$  memory, and  $O(1)$  time per item.
  - **Quality?**  
What is the probability that  $x_i \in X$  at some time  $t \geq i$ ?

# Reservoir Sampling II: Quality

- At any time step  $t$ , for  $i \leq t$ ,  $Pr[x_i \in X] = k/t$

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

**Reservoir sampling**

Sampling in a sliding  
window

# Reservoir Sampling II: Quality

- At any time step  $t$ , for  $i \leq t$ ,  $Pr[x_i \in X] = k/t$
- The proof is by induction on  $t$ .

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Reservoir Sampling II: Quality

- At any time step  $t$ , for  $i \leq t$ ,  $Pr[x_i \in X] = k/t$
- The proof is by induction on  $t$ .
  - Base  $t = k$ :  $Pr[x_i \in X] = 1$ , for  $i = 1, \dots, k$ .
  - Induction hypothesis: true for time steps up to  $t - 1$ 
    - $Pr[x_t \in X] = k/t$
    - For  $i < t$ ,  $x_i \in X$  when  $x_t$  is not selected and  $x_i$  was in the sample at step  $t - 1$ , or when  $x_t$  is selected,  $x_i$  was in the sample at step  $t - 1$  and  $x_i$  is not evicted.

$$\begin{aligned} Pr[x_i \in X] &= \left(1 - \frac{k}{t}\right) \frac{k}{t-1} + \frac{k}{t} \frac{k}{t-1} \left(1 - \frac{1}{k}\right) \\ &= \frac{k}{t-1} - \frac{k}{t} \frac{k}{t-1} \frac{1}{k} = \frac{k}{t-1} - \frac{1}{t} \frac{k}{t-1} = \frac{k}{t} \end{aligned}$$

# Reservoir Sampling for Sliding Windows

- **Problem:** Maintain a uniform sample of  $k$  items from the last  $w$  items.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Reservoir Sampling for Sliding Windows

- **Problem:** Maintain a uniform sample of  $k$  items from the last  $w$  items.
- **Why reservoir sampling does not work?**

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Reservoir Sampling for Sliding Windows

- **Problem:** Maintain a uniform sample of  $k$  items from the last  $w$  items.
- **Why reservoir sampling does not work?**
  - Suppose an element in the reservoir expires
  - Need to replace it with a randomly-chosen element from the current window
  - But, **we have no access to past data!**

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Reservoir Sampling for Sliding Windows

- **Problem:** Maintain a uniform sample of  $k$  items from the last  $w$  items.
- **Why reservoir sampling does not work?**
  - Suppose an element in the reservoir expires
  - Need to replace it with a randomly-chosen element from the current window
  - But, **we have no access to past data!**
  - Could store the entire window but this would require  $O(w)$  memory.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window



# Chain-Sampling for Sliding Windows

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

**Sampling in a sliding  
window**

# Chain-Sampling for Sliding Windows

## ■ Algorithm: ( $k = 1$ )

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Chain-Sampling for Sliding Windows

## ■ Algorithm: ( $k = 1$ )

- Maintain a reservoir sample for the first  $w$  items in  $s$ , but
- whenever an element  $x_i$  is selected, choose an index  $j \in [w]$  uniformly at random,  $x_{i+j}$  will be the replacement for  $x_i$ .
- For  $t > w$ , when  $t = i + j$ , set  $x = x_{i+j}$  (and choose the next replacement).

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Chain-Sampling for Sliding Windows

## ■ Algorithm: ( $k = 1$ )

- Maintain a reservoir sample for the first  $w$  items in  $s$ , but
- whenever an element  $x_i$  is selected, choose an index  $j \in [w]$  uniformly at random,  $x_{i+j}$  will be the replacement for  $x_i$ .
- For  $t > w$ , when  $t = i + j$ , set  $x = x_{i+j}$  (and choose the next replacement).

## ■ Analysis

- 1 pass,  $O(\log n + \log w)$  space and  $O(1)$  time per item.
- Provides a uniform sample.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Chain-Sampling for Sliding Windows

## ■ Algorithm: ( $k = 1$ )

- Maintain a reservoir sample for the first  $w$  items in  $s$ , but
- whenever an element  $x_i$  is selected, choose an index  $j \in [w]$  uniformly at random,  $x_{i+j}$  will be the replacement for  $x_i$ .
- For  $t > w$ , when  $t = i + j$ , set  $x = x_{i+j}$  (and choose the next replacement).

## ■ Analysis

- 1 pass,  $O(\log n + \log w)$  space and  $O(1)$  time per item.
- Provides a uniform sample.
- For higher values of  $k$ , run  $k$  parallel chain samples.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Chain-Sampling for Sliding Windows

## ■ Algorithm: ( $k = 1$ )

- Maintain a reservoir sample for the first  $w$  items in  $s$ , but
- whenever an element  $x_i$  is selected, choose an index  $j \in [w]$  uniformly at random,  $x_{i+j}$  will be the replacement for  $x_i$ .
- For  $t > w$ , when  $t = i + j$ , set  $x = x_{i+j}$  (and choose the next replacement).

## ■ Analysis

- 1 pass,  $O(\log n + \log w)$  space and  $O(1)$  time per item.
- Provides a uniform sample.
- For higher values of  $k$ , run  $k$  parallel chain samples. With high probability, for large enough  $w$ , such chains will not intersect.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Chain-Sampling: total time

Data stream  
models

Graph  
streams

Connectedness

Maximum matching

Sampling

Reservoir sampling

**Sampling in a sliding  
window**

# Chain-Sampling: total time

- $k = 1$
- The algorithm perform changes only along the chain of selected items.

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window



# Chain-Sampling: total time

- $k = 1$
- The algorithm perform changes only along the chain of selected items.
- The number of possible chains of elements with more than  $x$  data elements is bounded by the number of partitions of  $m$  into  $x$  ordered integer parts, which is bounded by  $\binom{m}{x}$ .
- Each such chain has probability at most  $m^{-x}$ .

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling  
Sampling in a sliding  
window

# Chain-Sampling: total time

- $k = 1$
- The algorithm perform changes only along the chain of selected items.
- The number of possible chains of elements with more than  $x$  data elements is bounded by the number of partitions of  $m$  into  $x$  ordered integer parts, which is bounded by  $\binom{m}{x}$ .
- Each such chain has probability at most  $m^{-x}$ .
- The probability of updating  $x$  steps is therefore at most  $\binom{m}{x} m^{-x}$ .
- Using Stirling's approximation we get the bound  $\left(\frac{e}{x}\right)^x$ .

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Chain-Sampling: total time

- $k = 1$
- The algorithm perform changes only along the chain of selected items.
- The number of possible chains of elements with more than  $x$  data elements is bounded by the number of partitions of  $m$  into  $x$  ordered integer parts, which is bounded by  $\binom{m}{x}$ .
- Each such chain has probability at most  $m^{-x}$ .
- The probability of updating  $x$  steps is therefore at most  $\binom{m}{x} m^{-x}$ .
- Using Stirling's approximation we get the bound  $\left(\frac{e}{x}\right)^x$ .
- For  $x = O(\log m)$  this is less than  $m^{-c}$ , for some constant  $c$

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window

# Chain-Sampling: total time

- $k = 1$
- The algorithm perform changes only along the chain of selected items.
- The number of possible chains of elements with more than  $x$  data elements is bounded by the number of partitions of  $m$  into  $x$  ordered integer parts, which is bounded by  $\binom{m}{x}$ .
- Each such chain has probability at most  $m^{-x}$ .
- The probability of updating  $x$  steps is therefore at most  $\binom{m}{x} m^{-x}$ .
- Using Stirling's approximation we get the bound  $\left(\frac{e}{x}\right)^x$ .
- For  $x = O(\log m)$  this is less than  $m^{-c}$ , for some constant  $c$
- With high probability the number of updates is  $O(\log m)$ .

Data stream  
models

Graph  
streams

Connectedness  
Maximum matching

Sampling

Reservoir sampling

Sampling in a sliding  
window